

Vishwanth Marri

+1 513-348-0777 | marrivishwanth24@gmail.com | [LinkedIn](#)  | [GitHub](#)  | [Portfolio](#)  | Dallas, Texas

SUMMARY

Full-stack AI engineer with production experience building LLM agent systems, RAG pipelines, evaluation frameworks, and agentic workflows using the Claude API, LangChain, MCP, and pgvector. Ships end-to-end features across React, Next.js, TypeScript, Node.js, Python, FastAPI, and AWS. Daily user of Claude Code and Cursor – using AI tools to build faster without sacrificing correctness.

TECHNICAL SKILLS

Programming Languages: Python, JavaScript, TypeScript, Java, SQL, C++

Web Technologies & Frameworks: React, Next.js, Node.js, Express, FastAPI, Flask, HTML5/CSS3

AI & ML: Claude API, LangChain, RAG Pipelines, MCP, Agent Orchestration, Prompt Engineering, Evaluation Frameworks, Scikit-learn

Databases & Caching: PostgreSQL, MySQL, MongoDB, Redis, pgvector, Elasticsearch

Cloud & DevOps: AWS (EC2, S3, RDS, Lambda), Docker, Kubernetes, GitHub Actions, CI/CD

AI Dev Tools: Claude Code, Cursor, GitHub Copilot

Core Concepts: REST APIs, Microservices, OAuth2/JWT, RBAC, System Design, Agile/Scrum

PROFESSIONAL EXPERIENCE

Kontaktsource

Aug 2025 – July 2026

Software Engineer Intern

Miami, Florida

- Engineered full HubSpot CRM integration layer handling OAuth2 auth flows, rate limiting, webhook processing, and undocumented API edge cases, automating onboarding workflows and reducing data-sync errors by 90%
- Architected backend automation layer using Node.js and FastAPI to streamline membership onboarding workflows, eliminating repetitive data processing steps and improving platform reliability end-to-end
- Accelerated frontend performance by refactoring React and TypeScript components with lazy loading and code splitting, reducing page load times by 40% across all devices
- Secured platform by implementing OAuth2 authorization, JWT validation, and adaptive rate-limiting at the API layer, reducing fraudulent sign-ups to near zero

Infotact Solutions

June 2022 – May 2024

Full Stack Developer

Hyderabad, India

- Architected and delivered full-stack enterprise applications using the MERN stack, automating complex business workflows and reducing manual processing time by 40%
- Containerized microservices with Docker and deployed on AWS EC2 via GitHub Actions CI/CD pipelines, cutting release cycles by 35% and improving deployment consistency
- Designed and maintained RESTful APIs in Node.js and Express with OAuth2 authentication, rate limiting, and webhook handling, enabling reliable third-party integrations across multi-tenant client deployments
- Built unit and integration test suites achieving 92% coverage across the full MERN stack, enforcing quality through structured code reviews and Agile sprint ceremonies
- Optimized PostgreSQL and MongoDB query performance through indexing and Redis caching, reducing average API response times by 30%

PROJECTS

AI-Powered Document Q&A System | *FastAPI, pgvector, LangChain, Claude API, MCP, TypeScript, React* | [Live Demo](#)

- Designed and deployed a production RAG system end-to-end: document ingestion pipeline, pgvector embedding and indexing, LangChain agent orchestration with MCP and tool use, context engineering, structured outputs, and real-time Claude API streaming with source citations – deployed on AWS EC2
- Diagnosed and resolved a race condition causing silent write collisions under concurrent user load by tracing failures through structured logs and fixing with an async job queue, maintaining sub-200ms retrieval latency
- Built evaluation framework to catch retrieval failures and hallucinations before they reached users; implemented multi-turn memory patterns for coherent long-session reasoning without redundant retrieval

Product Review Sentiment Analysis Pipeline | *Python, FastAPI, scikit-learn, TF-IDF, REST, pytest* | [Live Demo](#)

- Built and deployed a FastAPI microservice classifying product reviews as positive/negative/neutral with a TF-IDF + Logistic Regression model, serving real-time single and batch inference over REST with per-prediction interpretability
- Achieved 81% accuracy / 0.81 weighted-F1 (3-class, 5-fold cross-validation) through a fully reproducible; hardened with a 13-test pytest suite and pinned dependencies so the committed model deploys identically to production.

EDUCATION

University of Cincinnati

Cincinnati, OH

Masters of Engineering in Computer Science CGPA 3.9/4

Aug 2024 – April 2026